PromoScout: Automated Promotional Clip Discovery & Ranking

Candidate Solution Document — Kuku AI Generalist Assignment (Problem 1: Find the Promos / Ads) | Target: *I Married a Broken Soldier. He was a Duke* (Ep 1–50)

---

1. Problem Choice & Framing

Microdramas live or die by performance marketing on TikTok, Instagram Reels, and YouTube Shorts. Each 1–2 minute vertical episode is packed with dramatic escalation, but editors currently spend hours manually scrubbing through footage to find 20–60-second promo clips. With weekly releases, this editorial process is an unsustainable bottleneck.

Scoping & Sensible Assumptions (Addressing Brief Silence):

- **The Promo Hook (First 3 Seconds):** Scrolling users decide within 3 seconds whether to stay. Clips must open on immediate narrative ignition—a slap, public humiliation, or shocking reveal—at the screenplay's timestamp resolution (0.01 s).
- **The Cliffhanger (Unresolved Tension):** The clip must cut precisely during peak unresolved tension (an unanswered ultimatum or gasp) rather than letting the scene resolve.
- **Standalone Comprehension (Cold Read):** Viewers seeing the ad have never seen the series. The conflict must be self-evident without needing external lore or genealogies.
- **Screenplay Confidentiality:** Screenplay texts are proprietary. Prompts, raw outputs, and data files are excluded from public git; raw outputs are shared via a link-only Google Drive folder.

2. Approach, Architecture & Scale

Core Architecture Principle: *"Python owns timestamps, boundaries, and deterministic math; LLMs own creative judgment, candidate discovery, and comparative ranking."*

Component	Deterministic Python Engine	Creative LLM Stages (Groq GPT-OSS)
Parsing & Bounds	Regex parses 3,727 events (0 misses). Enforces strict 20.0s–60.0s boundaries to 0.01s resolution.	N/A ($0 token spend on deterministic arithmetic).
Candidate Discovery	Tracks event ordering; single feedback retry if proposed clip is under 20s.	**Groq openai/gpt-oss-120b:** Proposes 1–3 clips per episode with verbatim quotes and spoiler risk.
Duration Repair & Validation	Trims trailing non-dialogue buffer if duration is 60–66s. Strictly verifies quote matches, hook in first 3s, cliffhanger in final 5s.	Prompted with duration feedback on short clips rather than arbitrary programmatic expansion.
Evaluation & Ranking	Deduplicates intra-episode candidates (IoU > 0.5). Normalizes 5-metric formula. Enforces max 2 promos per episode across Top 10.	**Evaluator (gpt-oss-20b):** Anonymized cold-read ([SPEAKER A/B], rel 0.0s, no camera tags). **Ranker (gpt-oss-120b):** Shuffled listwise ranking.

Scalability, Rate Limiting & Token Economics:

- **Sliding-Window Throttling (--rpm):** Enforces requests-per-minute limits (default 15 RPM for provider free tiers) to prevent 429 throttling.
- **Deterministic Disk Caching:** Cache keys hashed with *(model + prompt + temperature + run_id)*. Passing separate run-IDs isolates runs for genuine stability tests.
- **Token Economics:** Processing 50 episodes requires ~100 LLM calls (~140k input tokens, ~15k output tokens), costing pennies and fitting within free-tier limits.
- **Human-in-the-Loop:** Editors inspect the standalone interactive HTML dashboard to approve or reject ranked clips, feeding decisions back into scoring weights.

3. Quality Checking, Benchmarks & Honest Failures

Rather than relying on unverified virality claims, PromoScout is benchmarked with deterministic quantitative metrics:

Episode	Dramatic Arc & Key Moment	Acceptable Gold Window	System Candidate	IoU	Result
Ep 1	Altar betrayal: Juliana abandoned by Gerald	11.3s – 41.2s (29.9s)	1.0s – 23.5s (22.5s)	**0.918**	**Match**
Ep 13	Thorn bush search for the lost ring	5.4s – 35.0s (29.6s)	5.4s – 33.8s (28.4s)	**0.769**	**Match**
Ep 25	Confronting Juliana's injury & hidden pain	5.1s – 32.5s (27.4s)	5.1s – 30.5s (25.4s)	**0.639**	**Match**
Ep 38	Prison kindness vs bread threat	16.0s – 46.0s (30.0s)	74.5s – 87.9s (13.4s)	0.130	**Miss**
Ep 50	Forbidden secret meeting with Father Thomas	1.8s – 28.0s (26.2s)	1.8s – 27.5s (25.7s)	**0.829**	**Match**

Quantitative Metrics Summary:

- **HitRate@3 on Gold Set: 80.0% (4/5 episodes matched with IoU > 0.5).** Ground truth was annotated across an objective uniform sampling rule (Episodes 1, 13, 25, 38, 50) prior to evaluation.
- **Stability Test (Overlap@10): 80.0% consistency.** Evaluated across dual independent runs at temperature 0.3 using distinct run-IDs to prevent cache distortion.

Honest Failure Mode Analysis (Observed Failures):

1. **Overt Conflict Favored Over Subtle Tension (Ep 38 Miss):** Because the cold-read evaluator saw only dialogue text with zero context, it scored an aggressive verbal confrontation higher than an intimate, high-stakes scene with a sympathetic prison guard, missing subtle dramatic subtext. *Fix: Calibrate cold-read prompt to recognize quiet tension alongside explosive dialogue.*
2. **Buffer Trimming Softening Final Visual Reaction:** On clips exceeding 60s, algorithmic trimming of non-dialogue trailing buffer achieved duration compliance but occasionally clipped the final gasp reaction shot following a cliffhanger line. *Fix: Allow duration tolerance up to 62s if trailing event is a reaction camera shot.*
3. **Mid-Arc Chemical Hooks Penalized as Spoilers:** In Episode 25, romantic tension between Juliana and Alexander was flagged as a 'Medium Spoiler' (-5 penalty), even though displaying relationship chemistry is an ad hook rather than a plot ruin. *Fix: Differentiate premise chemistry hooks from true plot mystery spoilers.*

4. Prototype & Production Deliverables

- **CLI Orchestrator (run.py):** Modular CLI supporting Groq models (openai/gpt-oss-120b, openai/gpt-oss-20b), --evaluator-model, stage-by-stage execution (--stage, --resume, --inspect), --rpm, and --mock fixture replay.
- **Interactive Web Studio (app.py):** Streamlit UI with 5 tabs: Top Promos, Screenplay Explorer, Stage Audit & Anonymization Inspector, Benchmarks, and Direct Downloads.
- **Evaluation Suite (evaluate.py):** Recomputes HitRate@3 (80%) and genuine dual-run Overlap@10 (80%) with synthetic fixture detection.
- **Interactive Dashboard (outputs/report.html):** Standalone visual report with duration pills, 5-metric score breakdown, verbatim quotes, and filterable table.
- **Structured Exports & Raw Outputs:** outputs/final_rankings.json, final_rankings.csv, and unedited logs in raw_outputs/ (shared via link-only Google Drive folder).
- **3-Minute Video Walkthrough Script (docs/walkthrough_script.md):** Structured script covering problem, live run, report, and failure teardowns.